

ALLEANDRA ASYRAFIAN NURSANI – MID AI ENGINEER

Analytical database: Clickhouse (Scale plan, trial)

Editor: Visual Studio Code, MySQL Workbench

Language: MySQL, python(requests, pandas, mysql-connector-python, prettytables, loggings, matplotlib)

1. DATA PIPELINE

a. Extraction

- Retrieve daily production data from the production_logs

```
cursor.execute("""select
p.date,
sum(case when p.tons_extracted < 0 then 0 else p.tons_extracted end) as 'total_production_daily',
round(avg(p.quality_grade), 2) as 'average_quality_grade',
-- count(distinct p.mine_id) as 'total_production_daily',
round(count(case when m.operational_status = 'active' then 1 end) * 100/count(*),2) as 'equipment_utilization'
from production_logs p
left join mines m on m.mine_id = p.mine_id
group by p.date;
""")

rows = cursor.fetchall()
columns = [desc[0] for desc in cursor.description]
df_prod = pd.DataFrame(rows, columns=columns)

assert df_prod['equipment_utilization'].between(0, 100).all()

#print(df_prod.head(10))

#for row in cursor.fetchall():
#    print(row)

table = PrettyTable()
table.field_names = df_prod.columns.tolist()

for row in df_prod.head(4).values.tolist():
    table.add_row(row)

print(table)
```

date	total_production_daily	average_quality_grade	equipment_utilization
2024-07-01	1483.01	4.65	66.67
2024-07-02	1867.32	3.77	66.67
2024-07-03	1366.60	4.93	66.67
2024-07-04	2175.39	4.55	66.67

- Read the equipment_sensors.csv

```
#MERGE 1 (PROD TO EQUIPMENT)

import pandas as pd
from prettytable import PrettyTable

df_prod['date'] = pd.to_datetime(df_prod['date'])
df_equipment['timestamp'] = pd.to_datetime(df_equipment['timestamp'])

merge_1 = pd.merge(df_prod, df_equipment, left_on='date', right_on='timestamp', how='inner')

#print(merge_1.head(4))

#----#

table = PrettyTable()
table.field_names = merge_1.columns.tolist()

for row in merge_1.head(4).values.tolist():
    table.add_row(row)

# Print the table
print(table)
```

date	total_production_daily	average_quality_grade	equipment_utilization	timestamp	equipment_id	status	fuel_consumption	maintenance_alert
2024-07-01 00:00:00	1483.01	4.65	66.67	2024-07-01 00:00:00	TR001	maintenance	0.0	False
2024-07-01 00:00:00	1483.01	4.65	66.67	2024-07-01 00:00:00	TR002	idle	0.0	False
2024-07-01 00:00:00	1483.01	4.65	66.67	2024-07-01 00:00:00	TR003	maintenance	0.0	False
2024-07-01 00:00:00	1483.01	4.65	66.67	2024-07-01 00:00:00	TR004	idle	0.0	False

- Call weather API

```
#!/usr/bin/env python3
import requests
from prettytable import PrettyTable

url = "https://api.open-meteo.com/v1/forecast? \
'latitude=2.0167&longitude=117.3600' \
'daily-temperature_2m_mean,precipitation_sum' \
&timezone=Asia/Jakarta' \
&next_days=4' \
&start_date=2025-05-15' \
&end_date=2025-07-31'"

#print(requests.get(url).json())

df_weather = pd.DataFrame(requests.get(url).json()["daily"])

#print(df_weather.head(4))
#print(df[df["time"] == df["time"].max()])

table = PrettyTable()
table.field_names = df_weather.columns.tolist()

for row in df_weather.head(11).values.tolist():
    table.add_row(row)

# Print the table
print(table)
```

time	temperature_2m_mean	precipitation_sum
2025-05-15	26.0	7.2
2025-05-16	25.3	4.2
2025-05-17	25.9	9.6
2025-05-18	25.4	4.6
2025-05-19	26.8	10.9
2025-05-20	24.3	20.1
2025-05-21	24.8	5.3
2025-05-22	27.2	0.1
2025-05-23	27.7	6.7
2025-05-24	27.6	2.1
2025-05-25	28.1	1.8

b. Transformation

- **Total_production_daily**
- **Average_quality_grade**

```
#print(df_prod.head(10))

#for row in cursor.fetchall():
| # print(row)

table = PrettyTable()
table.field_names = df_prod.columns.to_list()

for row in df_prod.head(4).values.tolist():
| table.add_row(row)

print(table)

print(df_prod.dtypes)
```

✓ 0.0s

date	mine_id	total_production_daily	average_quality_grade
2024-07-01 00:00:00	1	235.62	3.95
2024-07-01 00:00:00	2	485.72	4.65
2024-07-01 00:00:00	3	761.67	5.35
2024-07-02 00:00:00	1	793.23	3.55

- **equipment_utilization:**

```
#-----
table = PrettyTable()
table.field_names = df_equipment.columns.to_list()

for row in df_equipment.head(11).values.tolist():
| table.add_row(row)

print(table)

print(df_equipment.dtypes)
```

✓ 0.1s

timestamp	equipment_id	status	fuel_consumption	maintenance_alert	date_only	equipment_utilization
2024-07-01 00:00:00	TR001	maintenance	0.0	False	2024-07-01 00:00:00	35.83
2024-07-01 00:00:00	TR002	idle	0.0	False	2024-07-01 00:00:00	35.83
2024-07-01 00:00:00	TR003	maintenance	0.0	False	2024-07-01 00:00:00	35.83
2024-07-01 00:00:00	TR004	idle	0.0	False	2024-07-01 00:00:00	35.83
2024-07-01 00:00:00	TR005	active	8.45	False	2024-07-01 00:00:00	35.83
2024-07-01 01:00:00	TR001	active	4.37	False	2024-07-01 00:00:00	35.83
2024-07-01 01:00:00	TR002	active	9.36	False	2024-07-01 00:00:00	35.83

- **fuel_efficiency**

```
total_coal = merge_1.drop_duplicates(subset='date_only')['total_production_daily'].sum()

#total fuel
total_fuel = merge_1.groupby(['date_only'])['fuel_consumption'].sum()

#avg fuel
fuel_per_ton = round((total_fuel / total_coal),7)*10000

merge_1['fuel_per_ton'] = merge_1["date_only"].map(fuel_per_ton)

print(merge_1.head(6))
#print(f"Average fuel consumption per ton: {fuel_per_ton:.4f}")
```

[268] ✓ 0.0s

date	mine_id	total_production_daily	average_quality_grade
0 2024-07-01	1	235.62	3.95
1 2024-07-01	1	235.62	3.95
2 2024-07-01	1	235.62	3.95
3 2024-07-01	1	235.62	3.95
4 2024-07-01	1	235.62	3.95
5 2024-07-01	2	485.72	4.65

timestamp equipment id status fuel consumption maintenance alert \

- **weather_impact**

```
merge_2['is_rainy'] = merge_2['precipitation_sum'] > 1.0

rainy_avg_current_month = merge_2[merge_2['is_rainy']].groupby(merge_2['time'].dt.to_period('M'))['total_production_daily'].mean()
non_rainy_avg_current_month = merge_2[~merge_2['is_rainy']].groupby(merge_2['time'].dt.to_period('M'))['total_production_daily'].mean()

#impact = rainy_avg - non_rainy_avg

merge_2['rainy_avg_current_month'] = merge_2['time'].dt.to_period('M').map(rainy_avg).round(2)
merge_2['non_rainy_avg_current_month'] = merge_2['time'].dt.to_period('M').map(non_rainy_avg).round(2)
merge_2 = merge_2.merge(impact.rename('impact'), on='time', how='left')

#print(f"Rainy day production vs. non-rainy: {impact:.2f} difference")
```

✓ 0.0s

c. handling

- avoid tons_extracted values that are negative

```
-- MAIN
select
p.date,
p.mine_id,
sum(case when p.tons_extracted < 0 then 0 else p.tons_extracted end) as 'total_production_daily',
round(avg(p.quality_grade), 2) as 'average_quality_grade'
-- count(distinct p.mine_id) as 'total_production_daily'
from production_logs p
left join mines m on m.mine_id = p.mine_id
group by p.date, p.mine_id;
```

- unknown flagging for missing sensor data

```
#EQUIPMENT
import pandas as pd
from prettytable import PrettyTable

df_equipment_dirty = pd.read_csv(r'C:\Users\allen\Desktop\SYNOPSIS CHALLENGE\synopsis ai engine')
df_equipment = df_equipment_dirty.map(lambda x: x if pd.notna(x) else "unknown")

df_equipment["timestamp"] = pd.to_datetime(df_equipment_dirty["timestamp"])
df_equipment["equipment_id"] = df_equipment_dirty["equipment_id"].astype('string')
df_equipment["status"] = df_equipment_dirty["status"].astype('string')
df_equipment["fuel_consumption"] = pd.to_numeric(df_equipment_dirty["fuel_consumption"])
df_equipment["maintenance_alert"] = df_equipment_dirty["maintenance_alert"].astype('bool')

df_equipment['date_only'] = df_equipment['timestamp'].dt.date
#df_equipment['date_only'] = pd.to_datetime(df_equipment['date_only'])
```

2. IMPLEMENT ETL SCRIPT

```
secure=True

#print("Result:", client.query("SELECT 1").result_set[0][0])
client.command("""
CREATE TABLE IF NOT EXISTS merged_prod_to_equipment(
    date DateTime,
    mine_id Int64,
    total_production_daily Float64,
    average_quality_grade Float64,
    timestamp DateTime,
    equipment_id String,
    status String,
    fuel_consumption Float64,
    maintenance_alert UInt8,
    date_only datetime,
    equipment_utilization Float64,
    fuel_per_ton Float64
) ENGINE = MergeTree()
ORDER BY date

client.insert_df('merged_prod_to_equipment', merge_1)
```

Python

Clickhouse

The screenshot shows the ClickHouse web interface. On the left, a sidebar contains navigation links like 'Getting started', 'SQL Console', 'Dashboards', 'Data sources', 'Backups', 'Settings', 'Monitoring', and 'Help'. The main area is divided into 'Tables' and 'Queries' tabs. Under 'Tables', a table named 'merged_prod_to_equipment' is listed with columns: date, mine_id, total_production_daily, average_quality_grade, timestamp, equipment_id, status, fuel_consumption, maintenance_alert, date_only, equipment_utilization, and fuel_per_ton. The 'Queries' tab is active, displaying a complex SQL query that involves creating a table, inserting data, and performing a series of joins and aggregations to calculate various metrics like total production, average quality grade, and fuel consumption per ton, grouped by date and mine_id.

3. VALIDATE DATA

a. Implement checks

- Total_production_daily non-negative
- equipment_utilization is between 0 and 100%.
- Confirm weather data is complete for each production day.

```
#Error log
import logging

logging.basicConfig(
    filename='log_error.txt',
    level=logging.ERROR,
    format='%(asctime)s - %(levelname)s - %(message)s'
)

#total_production_daily is not negative
if not(df_prod['total_production_daily']>= 0).all() :
    logging.error("'total_production_daily' values is below 0.")
else:
    print("✅ All total_production_daily values are non-negative.")

#equipment_utilization between 0%-100%.
if not(df_equipment['equipment_utilization'].between(0, 100).all()) :
    logging.error("'equipment_utilization' values not between 0 and 100.")
else:
    print("✅ All equipment_utilization percentage are within range.")

#weather data is complete for each production day
if df_weather.isnull().any(axis=1).any():
    missing_rows = df_weather[df_weather.isnull().any(axis=1)]
    logging.error(f"Incomplete weather data found on production days:\n{missing_rows}")
else:
    print("✅ All weather data entries are complete for each production day.")
```

[260] ✓ 0.0s

...
✅ All total_production_daily values are non-negative.
✅ All equipment_utilization percentage are within range.
✅ All weather data entries are complete for each production day.

b. Handle anomalies

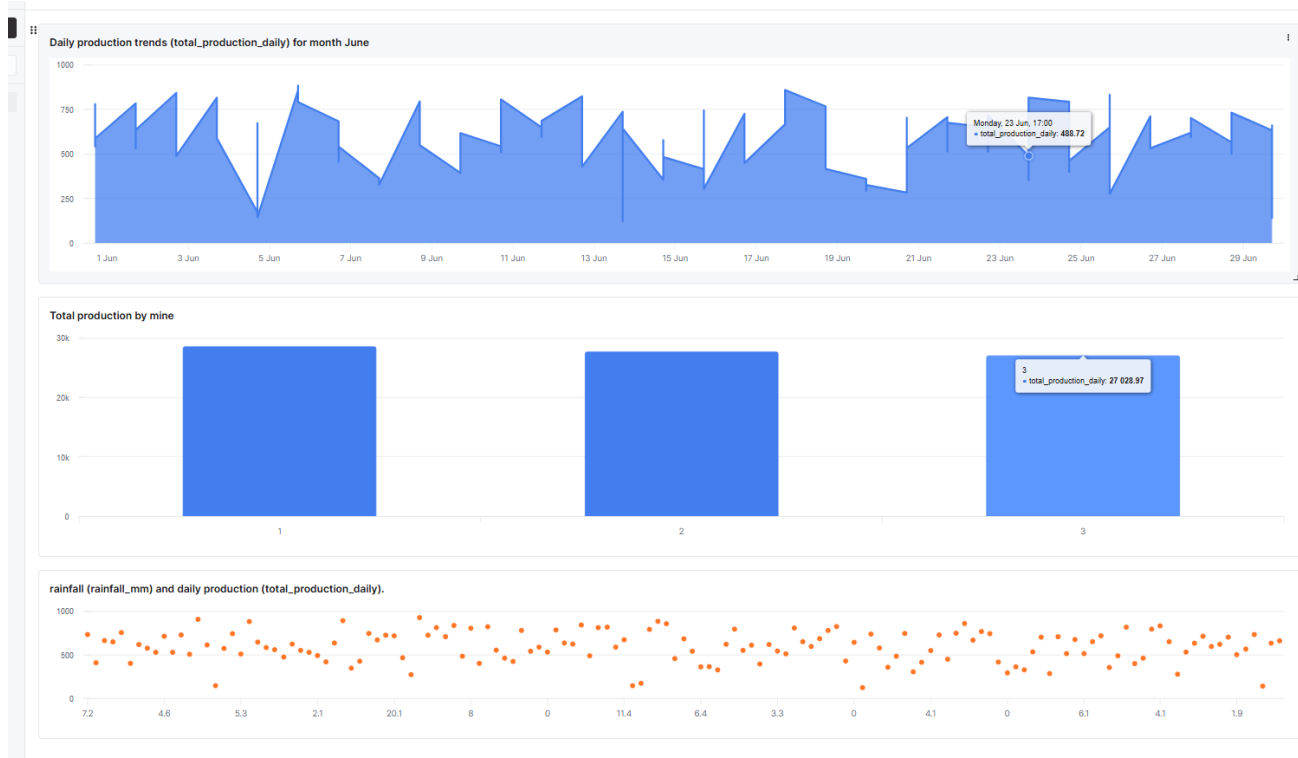
```
#Error log
import logging

logging.basicConfig(
    filename='log_error.txt',
    level=logging.ERROR,
    format='%(asctime)s - %(levelname)s - %(message)s'
)

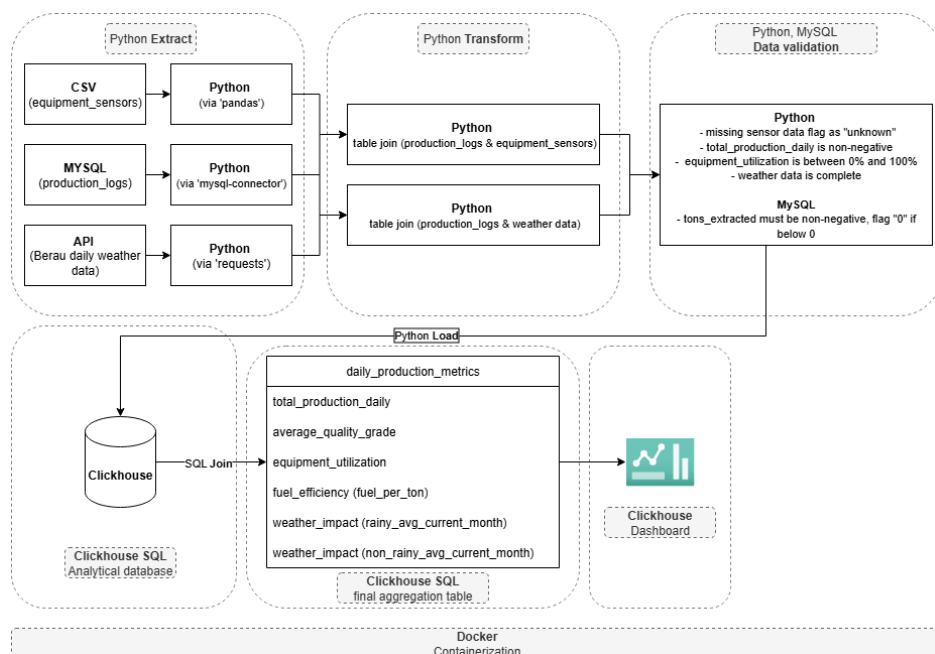
#total_production_daily is not negative
if not(df_prod['total_production_daily']>= 0).all() :
```

4. DASHBOARD

***note:** I'm most comfortable with Power BI, and can make great charts with PBI, however, because I used Clickhouse on an AWS VM (using Singapore server), I could not get the host url for PBI, so I ended up making basic charts in Clickhouse. I hope this is alright. Thank you!



5. DOCUMENT



6. PREDICTIVE MODEL GO TO NEXT PAGE

PREDCTIVE MODEL TO FORECAST THE NEXT 2 DAYS for mine_id = 1

```
import pandas as pd
import matplotlib.pyplot as plt
import statsmodels.api as sm

aggregated = merge_1.groupby(['mine_id', 'date'])['total_production_daily'].mean().re

# model for each mine_id
for mine in aggregated['mine_id'].unique():
    data = aggregated[aggregated['mine_id'] == mine].set_index('date')['total_product

    #arima
    model = sm.tsa.SARIMAX(data, order=(1, 0, 0), trend='c')
    results = model.fit()

    #forecast next 10 days
    forecast = results.get_forecast(steps=10)
    predicted_mean = forecast.predicted_mean
    conf_int = forecast.conf_int()

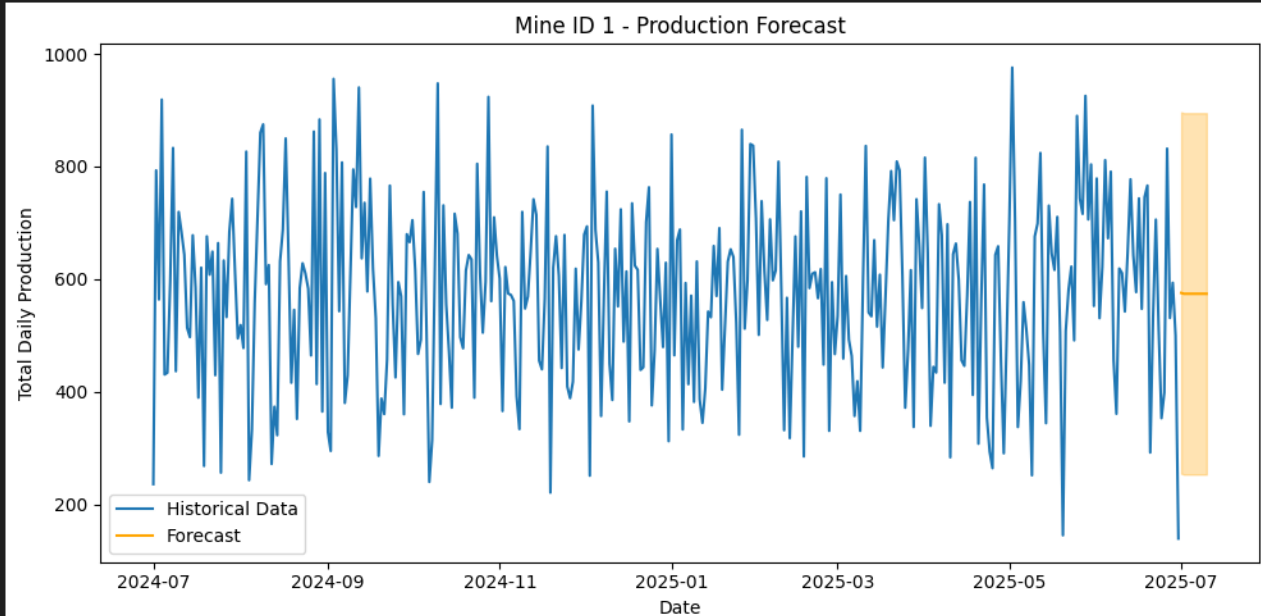
    # Create date range for forecast
    last_date = data.index[-1]
    forecast_dates = pd.date_range(last_date + pd.Timedelta(days=1), periods=10)

    # Plotting
    plt.figure(figsize=(10, 5))
    plt.plot(data, label='Historical Data')
    plt.plot(forecast_dates, predicted_mean, label='Forecast', color='orange')
    plt.fill_between(forecast_dates,
                    conf_int.iloc[:, 0],
                    conf_int.iloc[:, 1],
                    color='orange', alpha=0.3)

    plt.title(f"Mine ID {mine} - Production Forecast")
    plt.xlabel("Date")
    plt.ylabel("Total Daily Production")
    plt.legend()
    plt.tight_layout()
    plt.show()
```

✓ 62s

self._init_dates(dates, freq)



-- Thank you! --